



**CHALMERS**  
UNIVERSITY OF TECHNOLOGY



UNIVERSITY OF GOTHENBURG

# SQL, DDL, and Mapping E(E)R to Databases

LECTURE 7

**Dr. Philipp Leitner**



philipp.leitner@chalmers.se



@xLeitix




# LECTURE 7

**Covers ...**

**Parts of Chapter 6 and 7  
Chapter 9**

***Please read this up until next lecture!***





# What we will be covering

**We finally get started with some SQL**

**Defining tables, inserting data, altering tables**

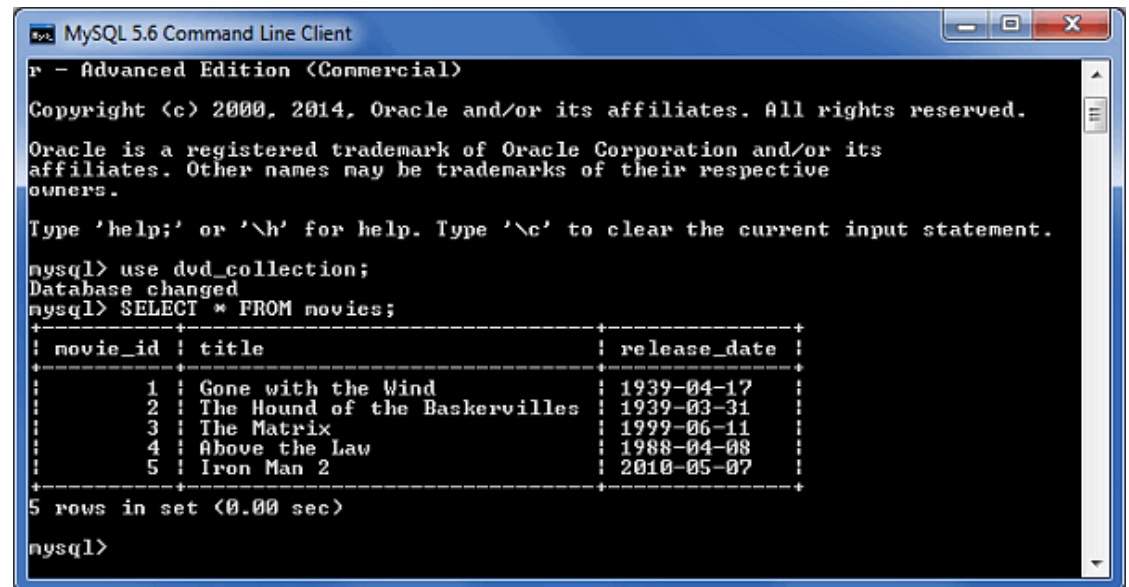
**Some patterns for mapping E(E)R to relations and databases**

# Reminder: DBMS

A database, and a DBMS, is basically a server process

SQL is the **interface** we largely use to interact with this

Conveniently, SQL follows RA **very closely**



```
MySQL 5.6 Command Line Client
r - Advanced Edition (Commercial)
Copyright (c) 2000, 2014, Oracle and/or its affiliates. All rights reserved.
Oracle is a registered trademark of Oracle Corporation and/or its
affiliates. Other names may be trademarks of their respective
owners.
Type 'help;' or '\h' for help. Type '\c' to clear the current input statement.
mysql> use dvd_collection;
Database changed
mysql> SELECT * FROM movies;
+-----+-----+-----+
| movie_id | title                               | release_date |
+-----+-----+-----+
| 1 | Gone with the Wind                  | 1939-04-17   |
| 2 | The Hound of the Baskervilles       | 1939-03-31   |
| 3 | The Matrix                          | 1999-06-11   |
| 4 | Above the Law                       | 1988-04-08   |
| 5 | Iron Man 2                          | 2010-05-07   |
+-----+-----+-----+
5 rows in set (0.00 sec)

mysql>
```

# SQL

SQL (pronounced “sequel” or sometimes “S-Q-L”)

Is a standardized **computer language** for interacting with relational databases

ISO/ANSI standard:

**Standardized** parts work for virtually all relational databases

Most have also some “non-standard” features (we won’t be covering them)



# Relational Database Systems

Open Source:

MySQL

**PostgreSQL**

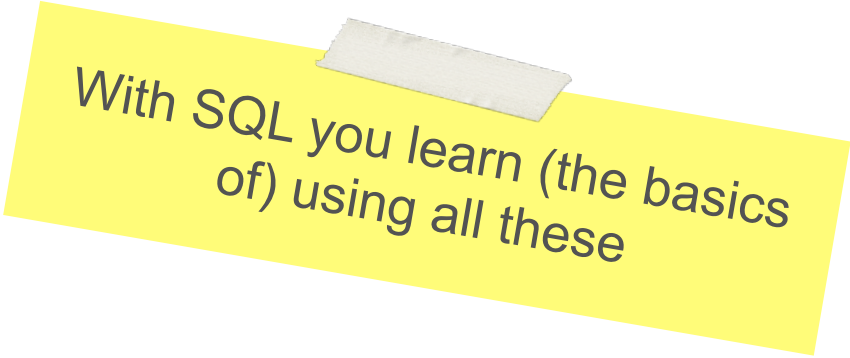
...

Commercial:

Oracle

SQL Server (Microsoft)

...



With SQL you learn (the basics  
of) using all these

# SQL in a Nutshell

Textual **domain-specific language**

A computer language like Java, but for **two** very specific purposes

Defining a database **schema**

DDL (database definition language)

Inserting, updating, and finding **data items**

DML (data manipulation language)

# Creating Tables

Basic DDL command is `CREATE TABLE`

Two versions:

`CREATE TABLE` (base relations)

Relation and its tuples are actually created and stored as a file by the DBMS

`CREATE VIEW` (create virtual relation)

Do not correspond to any physical file.



Let's do a small in-class example. Follow along on your computer if you can.

```
1. psql -U postgres (psql)
```

template1	postgres	UTF8	C	C	postgres=CtC/postgres
test	postgres	UTF8	C	C	=c/postgres
test2	postgres	UTF8	C	C	postgres=CtC/postgres

(5 rows)

```
postgres=# \c test2;
You are now connected to database "test2" as user "postgres".
test2=# drop database test;
test2=# ;
DROP DATABASE
test2=# drop database test2;
ERROR: cannot drop the currently open database
test2=# create database in_class;
CREATE DATABASE
test2=# \l

          List of databases
  Name      | Owner   | Encoding | Collate | Ctype  | Access privileges
-----
in_class    | postgres | UTF8     | C       | C      |
postgres    | postgres | UTF8     | C       | C      |
template0   | postgres | UTF8     | C       | C      | =c/postgres
template1   | postgres | UTF8     | C       | C      | =c/postgres
test2       | postgres | UTF8     | C       | C      | postgres=CtC/postgres
(5 rows)
```

```
test2=# drop database in_class;
DROP DATABASE
test2=# create database in_class;
CREATE DATABASE
test2=# \c in_class;
You are now connected to database "in_class" as user "postgres".
in_class=# \c test2;
You are now connected to database "test2" as user "postgres".
test2=# drop database in_class;
DROP DATABASE
test2=#
```

# Getting Started

- Log into your psql client (run `psql -U postgres` or something similar in a terminal)
- Type `\l` (show all databases)
- Create a database: type `create database in_class;`. Validate that the database is created.
- Connect to database: type `\c in_class;`

# Simple Example

```
CREATE TABLE student (  
    personnr CHAR(10),  
    name VARCHAR,  
    birthdate DATE,  
    PRIMARY KEY (personnr)  
);
```

# Validating that a table has been created

```
\dt;
```

(prints all tables in the database)

```
SELECT * FROM STUDENT;
```

(shows all content of a table, should currently be empty)

```
\d+ STUDENT;
```

(show table schema)



## More info

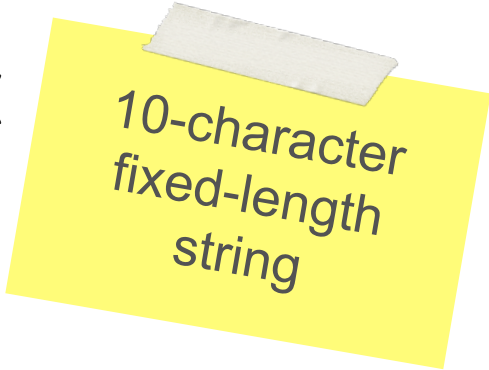


Table Name

```
CREATE TABLE STUDENT (  
    personnr CHAR(10),  
    name VARCHAR,  
    birthdate DATE,  
    PRIMARY KEY (personnr)  
);
```

## More info

```
CREATE TABLE STUDENT (  
    personnr CHAR(10),  
    name VARCHAR,  
    birthdate DATE,  
    PRIMARY KEY (personnr)  
);
```




10-character  
fixed-length  
string



## More info

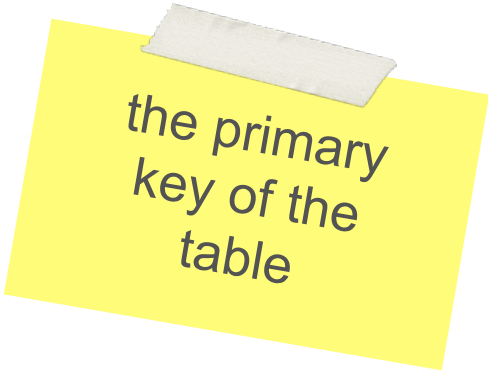
```
CREATE TABLE STUDENT (  
  personnr CHAR(10),  
  name VARCHAR,  
  birthdate DATE,  
  PRIMARY KEY (personnr)  
);
```



variable-length  
string

## More info

```
CREATE TABLE STUDENT (  
    personnr CHAR(10),  
    name VARCHAR,  
    birthdate DATE,  
    PRIMARY KEY (personnr)  
);
```



the primary  
key of the  
table



## Second Example

```
CREATE TABLE ENROLLMENT (  
    id INTEGER PRIMARY KEY,  
    coursename VARCHAR,  
    student CHAR(10),  
  
    FOREIGN KEY (student)  
        REFERENCES STUDENT(personnr)  
        ON DELETE CASCADE  
);
```



## More info

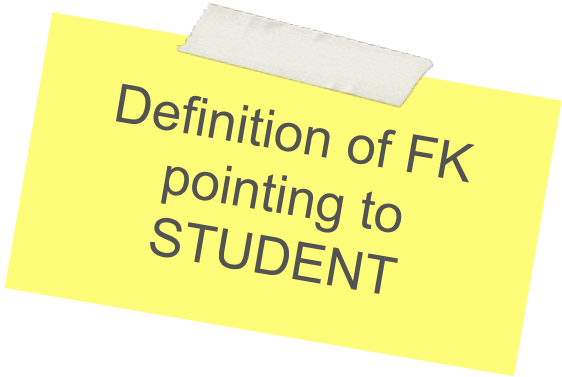
```
CREATE TABLE ENROLLMENT (  
    id INTEGER PRIMARY KEY,  
    coursenam VARCHAR,  
    student CHAR(10),  
  
    FOREIGN KEY (student)  
        REFERENCES STUDENT(personnr)  
        ON DELETE CASCADE  
);
```



Alternative way to  
specify PK

# More info

```
CREATE TABLE ENROLLMENT (  
    id INTEGER PRIMARY KEY,  
    coursename VARCHAR,  
    student CHAR(10),  
  
    FOREIGN KEY (student)  
        REFERENCES STUDENT(personnr)  
        ON DELETE CASCADE  
);
```

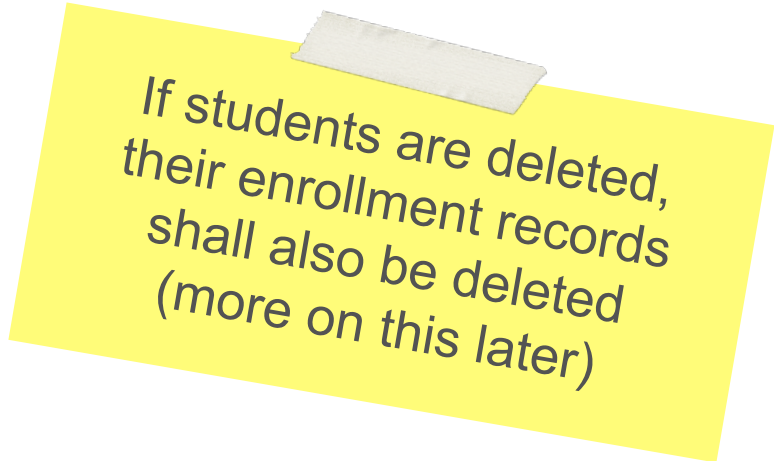


Definition of FK  
pointing to  
STUDENT



## More info

```
CREATE TABLE ENROLLMENT (  
    id INTEGER PRIMARY KEY,  
    coursename VARCHAR,  
    student CHAR(10),  
  
    FOREIGN KEY (student)  
        REFERENCES STUDENT(personnr)  
        ON DELETE CASCADE  
);
```



If students are deleted,  
their enrollment records  
shall also be deleted  
(more on this later)



### EMPLOYEE

Fname	Minit	Lname	<u>Ssn</u>	Bdate	Address	Sex	Salary	Super_ssn	Dno
John	B	Smith	123456789	1965-01-09	731 Fondren, Houston, TX	M	30000	333445555	5
Franklin	T	Wong	333445555	1955-12-08	638 Voss, Houston, TX	M	40000	888665555	5
Alicia	J	Zelaya	999887777	1968-01-19	3321 Castle, Spring, TX	F	25000	987654321	4
Jennifer	S	Wallace	987654321	1941-06-20	291 Berry, Bellaire, TX	F	43000	888665555	4
Ramesh	K	Narayan	666884444	1962-09-15	975 Fire Oak, Humble, TX	M	38000	333445555	5
Joyce	A	English	453453453	1972-07-31	5631 Rice, Houston, TX	F	25000	333445555	5
Ahmad	V	Jabbar	987987987	1969-03-29	980 Dallas, Houston, TX	M	25000	987654321	4
James	E	Borg	888665555	1937-11-10	450 Stone, Houston, TX	M	55000	NULL	1

### DEPARTMENT

Dname	<u>Dnumber</u>	Mgr_ssn	Mgr_start_date
Research	5	333445555	1988-05-22
Administration	4	987654321	1995-01-01
Headquarters	1	888665555	1981-06-19

### DEPT\_LOCATIONS

<u>Dnumber</u>	<u>Dlocation</u>
1	Houston
4	Stafford
5	Bellaire
5	Sugarland
5	Houston



CREATE TABLE  
statements to  
generate these  
tables

**CREATE TABLE EMPLOYEE**

( Fname	VARCHAR(15)	NOT NULL,
Minit	CHAR,	
Lname	VARCHAR(15)	NOT NULL,
Ssn	CHAR(9)	NOT NULL,
Bdate	DATE,	
Address	VARCHAR(30),	
Sex	CHAR,	
Salary	DECIMAL(10,2),	
Super_ssn	CHAR(9),	
Dno	INT	NOT NULL,

**PRIMARY KEY (Ssn),**

**CREATE TABLE DEPARTMENT**

( Dname	VARCHAR(15)	NOT NULL,
Dnumber	INT	NOT NULL,
Mgr_ssn	CHAR(9)	NOT NULL,
Mgr_start_date	DATE,	

**PRIMARY KEY (Dnumber),**

**UNIQUE (Dname),**

**FOREIGN KEY (Mgr\_ssn) REFERENCES EMPLOYEE(Ssn) );**

**CREATE TABLE DEPT\_LOCATIONS**

( Dnumber	INT	NOT NULL,
Dlocation	VARCHAR(15)	NOT NULL,

**PRIMARY KEY (Dnumber, Dlocation),**

**FOREIGN KEY (Dnumber) REFERENCES DEPARTMENT(Dnumber) );**

## 2 Basic Observations:

# Basic Structure

```
CREATE TABLE <TABLE_NAME> (  
    ATT1 DOM1 [NOT NULL] [CHECK (<STMT>)],  
    ATT2 DOM2 [NOT NULL],  
    ATT3 DOM3 [NOT NULL],  
    ...  
    PRIMARY KEY (ATTS),  
    [UNIQUE(ATTS), ]  
    [FOREIGN KEY(ATT3) REFERENCES TABLE2(ATTF)  
    [ON [DELETE,UPDATE]  
    [NO ACTION, RESTRICT, CASCADE, SET NULL, SET DEFAULT]]]  
);
```

(1) SQL *does not care* about capitalization of keywords

CREATE == create

(2) Statements end with ‘;’

# Attribute Data Types

## Basic data types

### Numeric data types

Integer numbers: `INTEGER` and `SMALLINT`

Floating-point (real) numbers: `FLOAT` and `DOUBLE PRECISION`

### Character-string data types

Fixed length: `CHAR ( n )`

Varying length: `VARCHAR ( n )`



# Attribute Data Types

## Bit-string data types

Fixed length: `BIT (n)`

Varying length: `BIT VARYING (n)`

## Boolean data type

Values of `TRUE` or `FALSE` or `NULL`

## DATE data type

Format `YYYY-MM-DD`

## BLOB / CLOB data type

[Binary | Character] Large Objects

E.g., binary or text files

# User-Defined Domains

## Example:

```
CREATE DOMAIN SSN_TYPE AS CHAR(8);
```

Useful if the same types are used in multiple places in the database  
Makes it easier to change them



# Constraints

## The fundamental ones:

Generic key constraint `UNIQUE ( ATTS )`

PK constraints `PRIMARY KEY ( ATTS )` (only one per table)

FK constraint

`FOREIGN KEY ( ATT3 ) REFERENCES TABLE2 ( ATTF )`



# Constraints

## Other constraints

DEFAULT <value>

NOT NULL

CHECK (implements arbitrary constraints)

Example:

```
Dnumber INT NOT NULL CHECK (  
    Dnumber > 0 AND Dnumber < 21);
```

# Constraints

## Multi-Attribute Primary Key

(has to be written after all attribute definitions)

Example:

```
PRIMARY KEY (essn, dependent_name);
```



# Constraints

Constraints can either be written

(1) **directly** when defining the attribute

```
Lname VARCHAR(20) NOT NULL
```

(2) or in a list **at the end** of the CREATE TABLE statement

```
CREATE TABLE EMPLOYEE(  
    Lname VARCHAR(20),  
    ...  
    NOT NULL(Lname)  
)
```

# Naming Constraints

Constraints can be explicitly named

May be useful to later comprehend them

```
CONSTRAINT LnameMustExist NOT NULL(Lname)
```

# Referential Integrity Triggers

For FOREIGN KEY we can tell the database what it should do with referential integrity violations

A primary key in a **referenced table** changes, what should happen with the **referencing foreign key**?

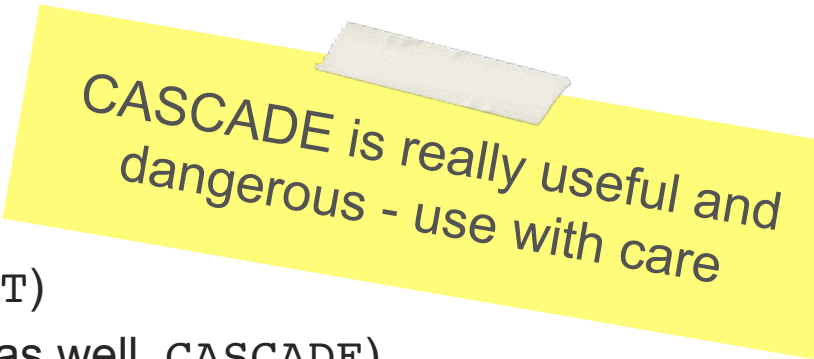
Can happen through updates or deletes

Four possibilities:

Reject update (default, NO ACTION, RESTRICT)

Cascade change (e.g., delete referencing rows as well, CASCADE)

Set to NULL or a different default (SET NULL, SET DEFAULT)



CASCADE is really useful and dangerous - use with care





# Changing Table Schemas

After a table is created, it can still be updated (`ALTER TABLE`) or deleted (`DROP TABLE`)

```
DROP TABLE TABLE1;
```

Deletes `TABLE1` and all data in the table

Cascades to other tables if configured

# Altering Schemas

Altering tables has a weird syntax. We do not focus on it in this course.

```
ALTER TABLE EMPLOYEE
```


```
    DROP COLUMN Address CASCADE
```

```
    ADD COLUMN NewAddress VARCHAR(30)
```

```
    ALTER COLUMN Fname SET DEFAULT 'Johan';
```


# In-Class Exercise

Try to create database tables that implement the following simple EER diagram:



# SQL DML - Data Manipulation Language

Create  
Read  
Update  
Delete



*Fundamental Operations Most  
Databases Will Support*

# Create, Updating, and Deleting

Three commands used to modify the database:

`INSERT`, `DELETE`, and `UPDATE`

`INSERT` inserts a new row into a table

`DELETE` removes one or more rows (that match a condition)

`UPDATE` changes one or more values in all rows that match a condition

We are going to cover the **read** part (querying) extensively in the next lecture.

# INSERT - two canonical forms

```
INSERT INTO EMPLOYEE VALUES (  
    'Hugo', 'Chavez', '324355423', '1962-12-30',  
    'Some Address', 'M', 37000, '325234523', 4  
);
```

Or:

```
INSERT INTO EMPLOYEE(Fname, Lname, Ssn) VALUES (  
    'Hugo', 'Chavez', '324323');
```

Not provided values in the second  
form become NULL

# INSERT - two canonical forms

```
INSERT INTO EMPLOYEE VALUES (  
    'Hugo', 'Chavez', '324355423', '1962-12-30',  
    'Some Address', 'M', 37000, '325234523', 4  
);
```

Or:

```
INSERT INTO EMPLOYEE(Fname, Lname, Ssn) VALUES (  
    'Hugo', 'Chavez', '324323');
```

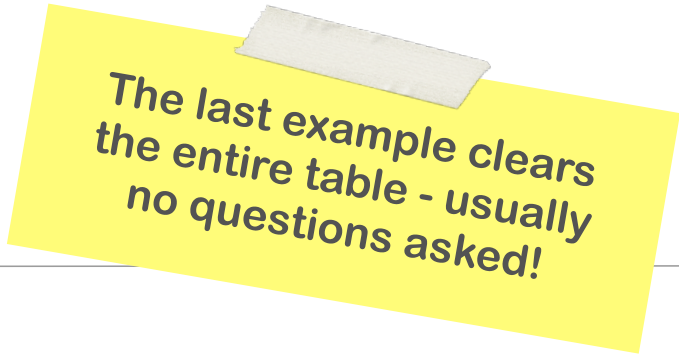
Watch out - only single backticks (!)

# DELETING data

```
DELETE FROM EMPLOYEE  
WHERE Lname = 'Chavez';
```

```
DELETE FROM EMPLOYEE  
WHERE Fname = 'Hugo';
```

```
DELETE FROM EMPLOYEE;
```



The last example clears  
the entire table - usually  
no questions asked!



# UPDATING data


```
UPDATE EMPLOYEE
```

```
SET Fname = 'Hugo', Lname = 'Chavez'
```

```
WHERE Ssn = '324323345';
```

```
UPDATE EMPLOYEE
```

```
SET Fname = 'Hugo', Lname = 'Chavez';
```



Everybody becomes  
Hugo Chavez - again, no  
questions asked.

# A word of warning

DBMSs are **expert tools** - they don't usually ask for confirmation, and there typically is no "undo" button.

**Think** before you execute a SQL statement, or make backups (e.g., database dump).

Many DBMSs return the number of **affected rows** for DML statements.

E.g., "17 rows deleted", "32453627 rows updated", ...

Allows sanity check whether command was correct

# How do get to tables from E(E)R

How do we get from our **logical domain model** in E(E)R to database tables?

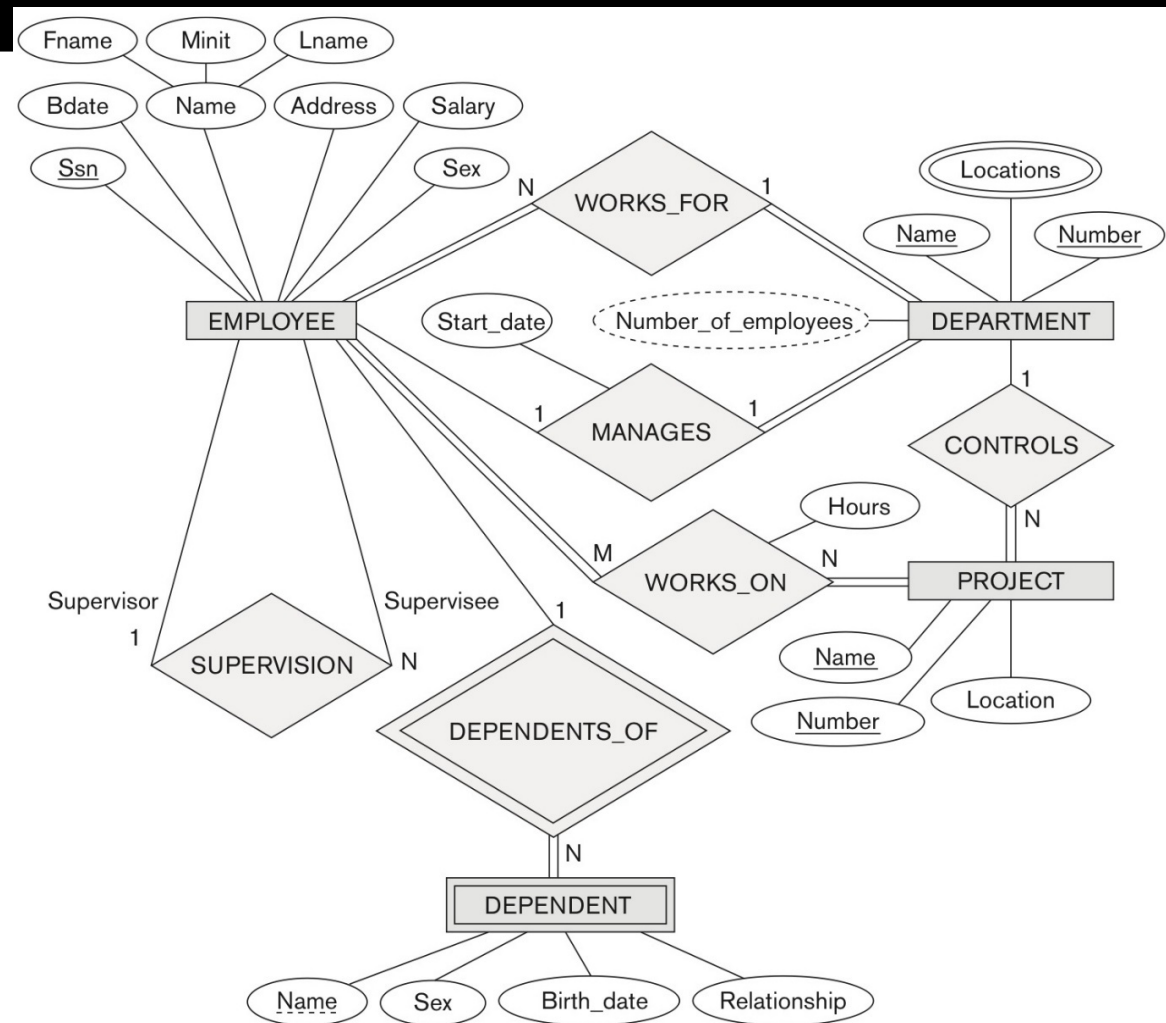
Mapping (mostly) straight-forward:

- Relations become tables

- Attributes become table attributes

- Attribute types and constraints need to be introduced

Let's discuss a  
few special cases.



# Mapping Relationships

## 1:N Relationship

Straight-forward to map via PK/FK (FK needs to be in the 'N' table)

## 1:1 Relationship

Map either like a 1:N relationship, *or*

Merge entities into single table (has all attributes from both relations)

## N:M Relationships need to be resolved via a cross-reference table

E.g., many EMPLOYEEs WORKS\_ON many PROJECTs

```
CREATE TABLE EMPLOYEE (...);  
CREATE TABLE PROJECT (...);  
CREATE TABLE EMP_PROJ(  
    Employee CHAR(9), Project CHAR(9),  
    PRIMARY KEY(Employee, Project),  
    FOREIGN KEY(Employee) REFERENCES EMPLOYEE(Ssn),  
    FOREIGN KEY(Project) REFERENCES PROJECT(Number)  
);
```

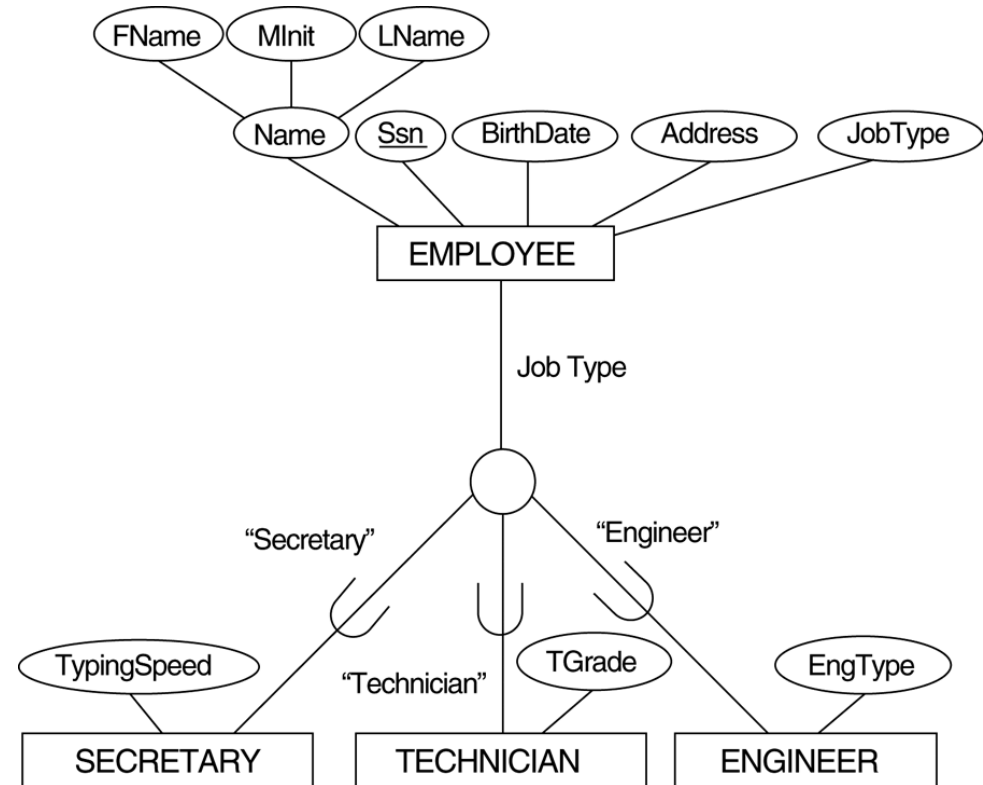
# Same basic approach can also be used for:

Mapping relationships with an **relationship attribute**

Mapping **composed attributes** (attribute becomes a separate table)

Mapping **multi-value attributes** (attribute again becomes a separate table)

# Mapping Inheritance





## Four different approaches, depending on type of inheritance.

Remember, EER inheritance comes in four flavors:

**Disjoint / total**

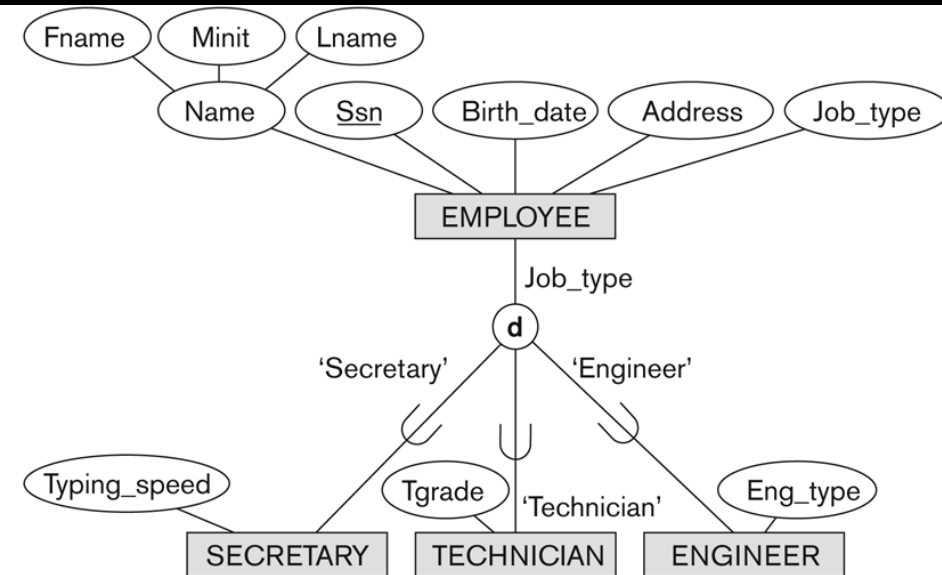
**Disjoint / partial**

**Overlapping / total**

**Overlapping / partial**

Correct mapping to tables depends on the flavor.

**Figure 4.4**  
EER diagram notation  
for an attribute-  
defined specialization  
on Job\_type.



## Disjoint / Partial

Map as four tables, one for the super-entity and one each for the sub-entities.

Use FKs to handle relationships.

(a) EMPLOYEE

<u>SSN</u>	FName	MInit	LName	BirthDate	Address	JobType
------------	-------	-------	-------	-----------	---------	---------

SECRETARY

<u>SSN</u>	TypingSpeed
------------	-------------


TECHNICIAN

<u>SSN</u>	TGrade
------------	--------

ENGINEER

<u>SSN</u>	EngType
------------	---------

# Disjoint / Total



Map as two tables:

(b) CAR

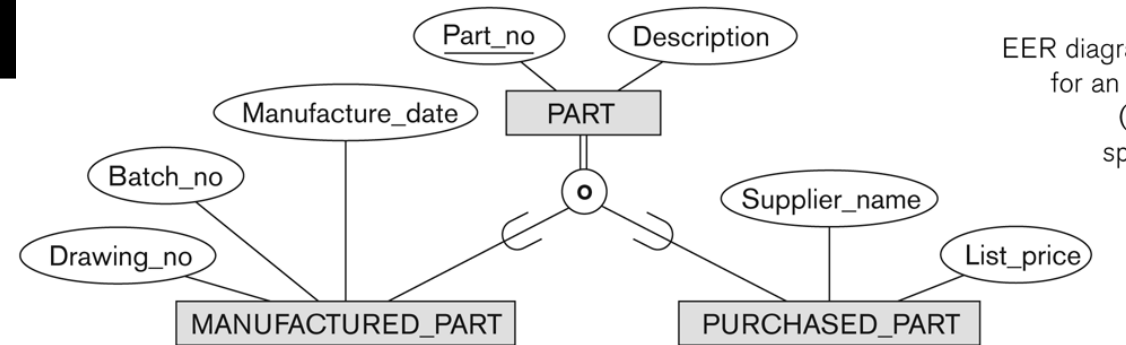
<u>VehicleId</u>	LicensePlateNo	Price	MaxSpeed	NoOfPassengers
------------------	----------------	-------	----------	----------------

TRUCK

<u>VehicleId</u>	LicensePlateNo	Price	NoOfAxles	
------------------	----------------	-------	-----------	--

Alternatively, map as three tables (VEHICLE, CAR, TRUCK) and join via FKs, just like for disjoint / partial.

# Overlapping / Total



Map as a single table with all attributes:

(d) PART

<u>PartNo</u>	Description	MFlag	DrawingNo	ManufactureDate	BatchNo	PFlag	SupplierName	ListPrice
---------------	-------------	-------	-----------	-----------------	---------	-------	--------------	-----------

All of the attributes coming from sub-entities need to be allowed to be NULL.

You don't necessarily need the sub-entity flags.

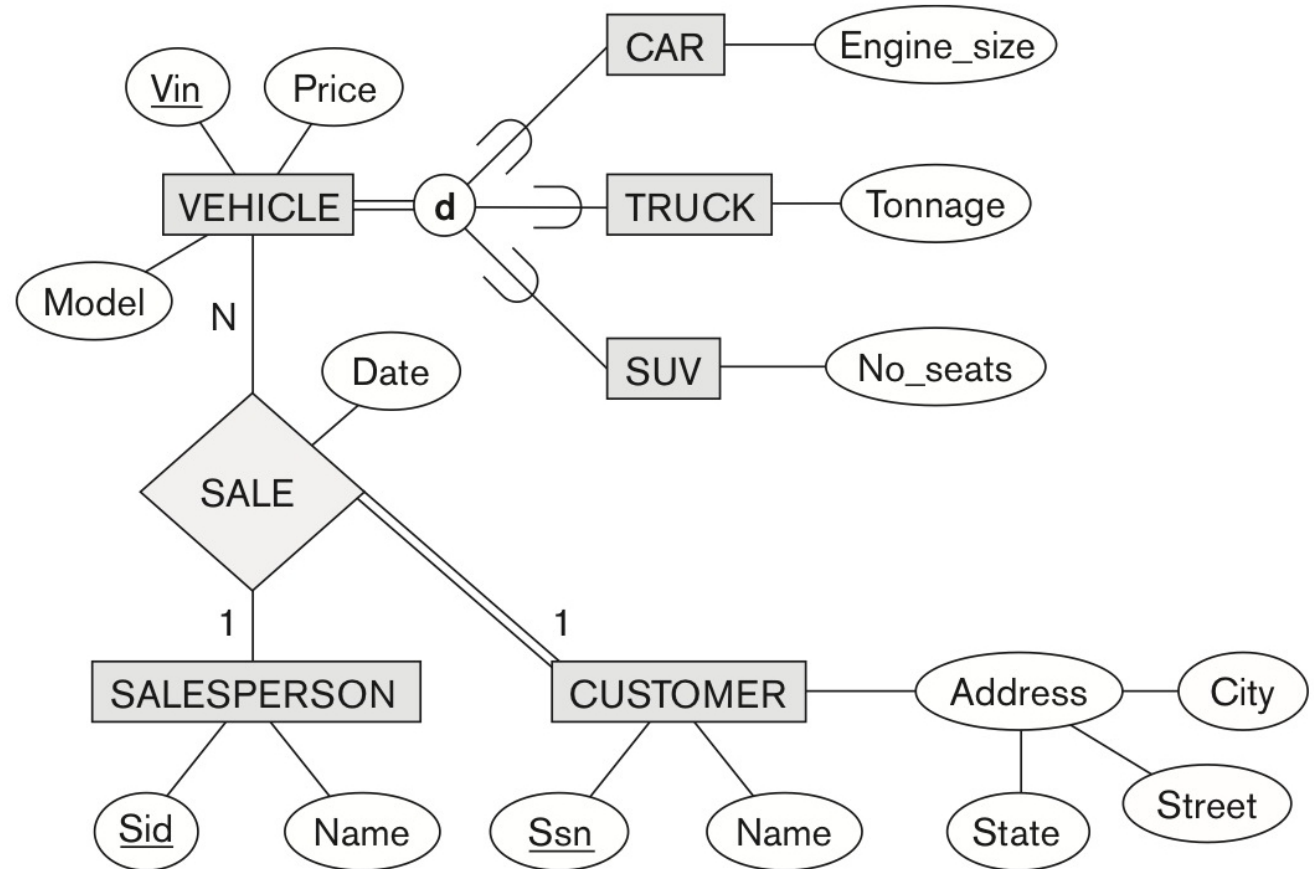


# Overlapping / Partial

Map the same as Overlapping / Total

Main difference is that for some rows, all attributes coming from sub-entities will be NULL.

## Another Exercise - Create the table model for this



# Key Takeaways

## **Basic DDL and DML SQL commands**

You need to train using them as part of the assignments!

## **Strategies for mapping E(E)R to SQL tables**

It's a fairly mechanical process, but it is important that you understand the general concepts