



**CHALMERS**  
UNIVERSITY OF TECHNOLOGY



UNIVERSITY OF GOTHENBURG

# Querying in SQL

LECTURE 8

**Dr. Philipp Leitner**



philipp.leitner@chalmers.se



@xLeitix



# Last lecture

**We started with SQL, and discussed how to create, drop, and alter tables, as well as how to insert and modify data.**

# Creating Tables

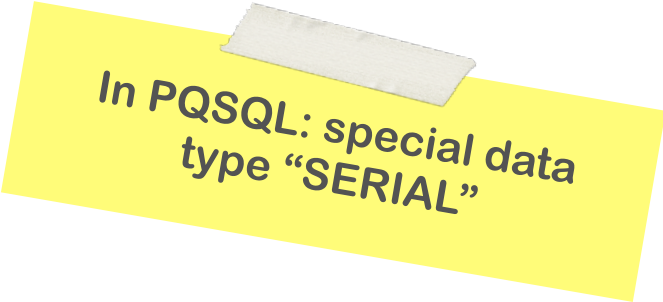
```
CREATE TABLE ENROLLMENT (  
    id INTEGER PRIMARY KEY,  
    coursename VARCHAR,  
    student CHAR(10),  
  
    FOREIGN KEY (student)  
        REFERENCES STUDENT(personnr)  
        ON DELETE CASCADE  
);
```



# Auto-Increment

(Most) relational databases have a feature to support **automatically generating** a primary key  
Usually through an automatically increasing ID counter

```
CREATE TABLE ENROLLMENT (  
    id SERIAL,  
    coursename VARCHAR,  
    student CHAR(10),  
  
    FOREIGN KEY (student)  
        REFERENCES STUDENT(personnr)  
        ON DELETE CASCADE  
);
```



In PQSQL: special data  
type "SERIAL"



# Constraints

UNIQUE ( ATTS )

PRIMARY KEY ( ATTS )

FOREIGN KEY ( ATT3 ) REFERENCES TABLE2 ( ATTF )

DEFAULT <value>

NOT NULL

CHECK <expression>

# Referential Integrity Triggers

For FOREIGN KEY we can tell the database what it should do with referential integrity violations

A primary key in a **referenced table** changes, what should happen with the **referencing foreign key**?

Can happen through updates or deletes

Four possibilities:

Reject update (default, NO ACTION, RESTRICT)

Cascade change (e.g., delete referencing rows as well, CASCADE)

Set to NULL or a different default (SET NULL, SET DEFAULT)

# INSERT

```
INSERT INTO EMPLOYEE VALUES (  
    'Hugo', 'Chavez', '324355423', '1962-12-30',  
    'Some Address', 'M', 37000, '325234523', 4  
);
```

Or:

```
INSERT INTO EMPLOYEE(Fname, Lname, Ssn) VALUES (  
    'Hugo', 'Chavez', '324323');
```

# DELETING data

```
DELETE FROM EMPLOYEE  
WHERE Lname = 'Chavez';
```

```
DELETE FROM EMPLOYEE  
WHERE Fname = 'Hugo';
```

```
DELETE FROM EMPLOYEE;
```



# UPDATING data


```
UPDATE EMPLOYEE
```

```
SET Fname = 'Hugo', Lname = 'Chavez'
```

```
WHERE Ssn = '324323345';
```

```
UPDATE EMPLOYEE
```

```
SET Fname = 'Hugo', Lname = 'Chavez';
```



Everybody becomes  
Hugo Chavez - again, no  
questions asked.

# Mapping EER to RM / SQL

General:

Entity types become tables, attributes become columns

One key becomes the primary key, all others `UNIQUE`

Data types need to be introduced


1:1 and 1:N relationships get mapped through PK/Oks

For N:M relationships and many special cases:


Use cross-reference table

# Cross-reference table

An N:M Relationship



With Cross-Reference Table



# Cross-reference table

E.g., many EMPLOYEEs WORKS\_ON many PROJECTs

```
CREATE TABLE EMPLOYEE (...);  
CREATE TABLE PROJECT (...);  
CREATE TABLE EMP_PROJ(  
    Employee CHAR(9), Project CHAR(9),  
    PRIMARY KEY(Employee, Project),  
    FOREIGN KEY(Employee) REFERENCES EMPLOYEE(Ssn),  
    FOREIGN KEY(Project) REFERENCES PROJECT(Number)  
);
```

# Mapping inheritance

**Disjoint / partial:** one table for each entity

**Disjoint / total:**

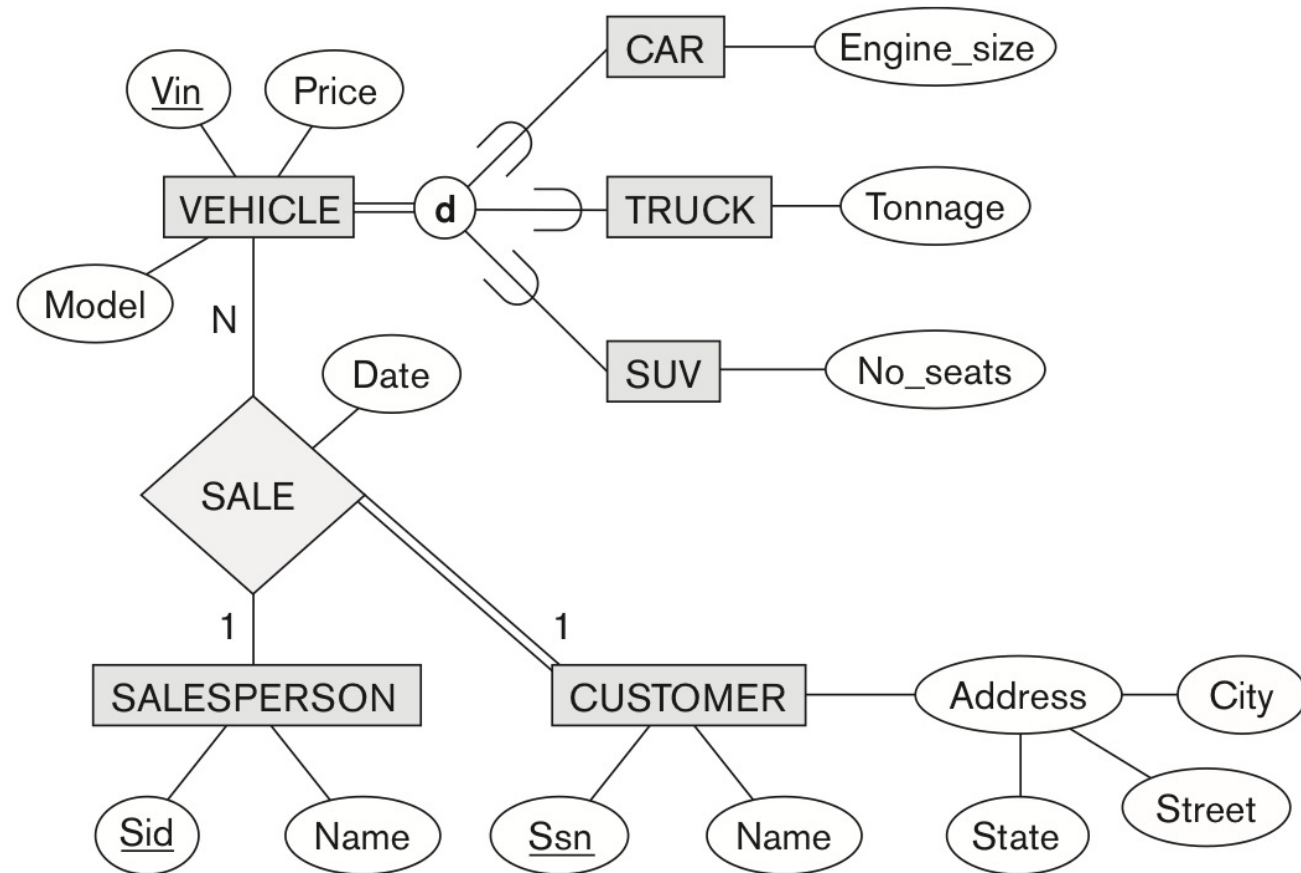
one table per concrete entity OR one table for each entity

**Overlapping / total:** single table with all attributes

**Overlapping / partial:** single table with all attributes

# In-Class Exercise

- (1) Produce a relational model for this EER diagram.
- (2) What SQL code would you use to create the tables?






# LECTURE 8

**Covers ...**

**Most of the remaining parts of  
Chapter 6 and 7**

***Please read this up until next lecture!***





# What we will be covering

## Writing SQL Queries


From the simply run-of-the-mill selects ...  
... over joins and aggregates ...  
... to subselects



# Basic Retrieval Queries (“READ”)

## SELECT statement

One basic statement for retrieving information from a database  
Idea follows relational algebra:



# Basic Structure

```
SELECT    <attribute list>  
FROM      <table list>  
WHERE     <condition>;
```

where

- <attribute list> is a list of attribute names whose values are to be retrieved by the query.
- <table list> is a list of the relation names required to process the query.
- <condition> is a conditional (Boolean) expression that identifies the tuples to be retrieved by the query.

# Basic Structure

```
SELECT <attribute list>  
FROM   <table list>  
WHERE  <condition>;
```

where

- <attribute list> is a list of attribute names whose values are to be retrieved by the query.
- <table list> is a list of the relation names required to process the query.
- <condition> is a conditional (Boolean) expression that identifies the tuples to be retrieved by the query.

PROJECT and RENAME

# Basic Structure

```
SELECT  <attribute list>  
FROM    <table list>  
WHERE   <condition>;
```

where

- <attribute list> is a list of attribute names whose values are to be retrieved by the query.
- <table list> is a list of the relation names required to process the query.
- <condition> is a conditional (Boolean) expression that identifies the tuples to be retrieved by the query.



SELECT



# Simple Example

**Query 0.** Retrieve the birth date and address of the employee(s) whose name is 'John B. Smith'.

```
Q0:  SELECT  Bdate, Address
      FROM    EMPLOYEE
      WHERE   Fname='John' AND Minit='B' AND Lname='Smith';
```

**EMPLOYEE**

Fname	Minit	Lname	<u>Ssn</u>	Bdate	Address	Sex	Salary	Super_ssn	Dno
John	B	Smith	123456789	1965-01-09	731 Fondren, Houston, TX	M	30000	333445555	5
Franklin	T	Wong	333445555	1955-12-08	638 Voss, Houston, TX	M	40000	888665555	5
Alicia	J	Zelaya	999887777	1968-01-19	3321 Castle, Spring, TX	F	25000	987654321	4
Jennifer	S	Wallace	987654321	1941-06-20	291 Berry, Bellaire, TX	F	43000	888665555	4
Ramesh	K	Narayan	666884444	1962-09-15	975 Fire Oak, Humble, TX	M	38000	333445555	5
Joyce	A	English	453453453	1972-07-31	5631 Rice, Houston, TX	F	25000	333445555	5
Ahmad	V	Jabbar	987987987	1969-03-29	980 Dallas, Houston, TX	M	25000	987654321	4
James	E	Borg	888665555	1937-11-10	450 Stone, Houston, TX	M	55000	NULL	1

**DEPARTMENT**

Dname	<u>Dnumber</u>	Mgr_ssn	Mgr_start_date
Research	5	333445555	1988-05-22
Administration	4	987654321	1995-01-01
Headquarters	1	888665555	1981-06-19

**DEPT\_LOCATIONS**

<u>Dnumber</u>	<u>Dlocation</u>
1	Houston
4	Stafford
5	Bellaire
5	Sugarland
5	Houston



# Simple Example

**Query 0.** Retrieve the birth date and address of the employee(s) whose name is 'John B. Smith'.

**Q0:**     **SELECT**     Bdate, Address  
         **FROM**       EMPLOYEE  
         **WHERE**     Fname='John' **AND** Minit='B' **AND** Lname='Smith';

<u>Bdate</u>	<u>Address</u>
1965-01-09	731 Fondren, Houston, TX

# Projection Wildcards

SQL supports the use of **wildcards** as part of projection (unlike RA)

```
SELECT * FROM EMPLOYEE;
```

(means “all attributes” from EMPLOYEE)



# Aliasing Tables

Tables can be given an **alias** during selection

Particularly important for joins

Mandatory if same table is selected more than one

```
SELECT E.Bdate FROM EMPLOYEE as E;
```



# Renaming Attributes

Two ways to rename attributes:

(1) in the attribute list:

```
SELECT Fname as Firstname FROM EMPLOYEE;
```

(2) as part of table selection

```
SELECT * FROM EMPLOYEE  
as E(Fn, Mi, Ln, Ssn, Bd, Addr, Sex, Sal, Sssn, Dno);
```

In some implementations you  
can drop the "as"



# Arithmetics

Arithmetic operations can be used in the SELECT and WHERE clause

```
SELECT Fname, Lname, 1.1 * Salary AS Increased_sal  
FROM EMPLOYEE;
```

(show what happened if everybody got a 10% salary increase)

# SELECTing Rows

Basic syntax: WHERE <boolean expression>

```
SELECT Fname as Firstname FROM EMPLOYEE  
WHERE Fname = 'John';
```

```
SELECT Fname FROM EMPLOYEE as E  
WHERE E.Fname = 'John';
```

```
SELECT Fname, Lname FROM EMPLOYEE  
WHERE Fname = 'John' AND Lname = 'Smith';
```

# Basic Operators for WHERE Clauses

Standard logical operators:

`=, <, <=, >, >=, <>` (note: `=`, not `==`)

Boolean operators:

`AND, OR`

Negation:

`NOT` (as in: `NOT ( a = 0 )` )

Is / Is Not Null:

`IS NULL, IS NOT NULL`



# Three-Value Logic

SQL uses a special kind of logic:  
TRUE, FALSE, UNKNOWN (NULL)

**Table 7.1** Logical Connectives in Three-Valued Logic

(a)	<b>AND</b>	TRUE	FALSE	UNKNOWN
	TRUE	TRUE	FALSE	UNKNOWN
	FALSE	FALSE	FALSE	FALSE
	UNKNOWN	UNKNOWN	FALSE	UNKNOWN
(b)	<b>OR</b>	TRUE	FALSE	UNKNOWN
	TRUE	TRUE	TRUE	TRUE
	FALSE	TRUE	FALSE	UNKNOWN
	UNKNOWN	TRUE	UNKNOWN	UNKNOWN
(c)	<b>NOT</b>			
	TRUE	FALSE		
	FALSE	TRUE		
	UNKNOWN	UNKNOWN		



# Additional Operators

LIKE (for textual comparison):

`Fname LIKE '%ohn%'` (matches 'John', 'Ohn', 'Johns', ...)

The % is a wildcard for 0 or more other characters

`_` can be used to match exactly 1 other character

BETWEEN (to check if value is between two other values):

`Salary BETWEEN 30000 AND 40000`

IN (to check if value is in list):

`Fname IN('John', 'Herbert')`

# Basic Operators for WHERE Clauses

## First order logical operators:

IN e.g., Ssn IN (123456, 123457, 123458)

ALL e.g., Salary > ALL (<subquery>)

ANY e.g., Salary < ANY (<subquery>)

SOME (same as ANY)

EXISTS e.g., EXISTS(<subquery>)

UNIQUE e.g., UNIQUE(<subquery>)



We get to  
subqueries later



# Tables versus Relations

Remember: relations are mathematical sets

No duplicates, no ordering

Tables are ***not*** sets

Duplicates ok, ordered (typically in order of insertion)

# Duplicates in Tables

**All** SQL tables are allowed to have duplicates

Multiple rows where all attributes have the same values

In practice no large concern for “normal” tables because of keys

**However, the result of a query (which is also a table!) can and will often have duplicates.**



# Duplicates in Tables

Special SQL keyword to eliminate duplicates:

**Query 11.** Retrieve the salary of every employee (Q11) and all distinct salary values (Q11A).

**Q11:**    **SELECT**    **ALL** Salary  
          **FROM**     EMPLOYEE;

**Q11A:**   **SELECT**   **DISTINCT** Salary  
          **FROM**     EMPLOYEE;

You will virtually never type  
“all”, it’s the assumed default



# Ordering Query Results

Use **ORDER BY** clause: ORDER BY <list\_of\_attributes>

Keyword **DESC** to see result in a descending order of values

Keyword **ASC** (default) to specify ascending order explicitly

Placed at the **end** of a query

Multiple order attributes can be given, then **first has priority**, second used as tie breaker and so on

```
ORDER BY D.Dname DESC, E.Lname ASC, E.Fname ASC
```

# Joining Tables

Most queries are run against multiple tables (i.e., **join queries**)

**Query 1.** Retrieve the name and address of all employees who work for the 'Research' department.

```
Q1:  SELECT  Fname, Lname, Address
      FROM    EMPLOYEE, DEPARTMENT
      WHERE   Dname='Research' AND Dnumber=Dno;
```



**Query 1.** Retrieve the name and address of all employees who work for the 'Research' department.

**Q1:**     **SELECT**     Fname, Lname, Address  
         **FROM**       EMPLOYEE, DEPARTMENT  
         **WHERE**      Dname='Research' **AND** Dnumber=Dno;

<u>Fname</u>	<u>Lname</u>	<u>Address</u>
John	Smith	731 Fondren, Houston, TX
Franklin	Wong	638 Voss, Houston, TX
Ramesh	Narayan	975 Fire Oak, Humble, TX
Joyce	English	5631 Rice, Houston, TX

# Basic Syntax of a JOIN

Most simple syntax is just listing multiple tables in FROM clause:

```
FROM Employee, Department
```

Note that this is equivalent to the **Cartesian Product  $A \times B$** , which we said is usually not very useful

Usually we combine the JOIN with **explicitly linking** the tables through their PK/FKs

```
SELECT E.Fname, E.Lname, D.Dname  
FROM Employee as E, Department as D  
WHERE E.Dno = D.Dnumber;
```

# Explicit Linking via On Clauses

Instead of establishing links in the WHERE clause we can also use an explicit “ON”

```
SELECT E.Fname, E.Lname, D.Dname  
FROM Employee as E INNER JOIN Department as D  
    ON E.Dno = D.Dnumber;
```

A bit more verbose, but makes clearer what is going on  
Also more consistent with OUTER JOIN syntax (see later)



# Self-Joins

Counter-intuitively, a table can also be joined with itself:

```
SELECT  E.Fname, E.Lname, S.Fname, S.Lname  
FROM    EMPLOYEE AS E, EMPLOYEE AS S  
WHERE   E.Super_ssn=S.Ssn;
```

In this case aliasing the tables is **necessary**.

# Basic Syntax of a JOIN

Complex join queries can get fairly verbose

```
SELECT *  
FROM A as a, B as b, C as c, D as d  
WHERE a.pk = b.fk AND b.pk = c.fk AND c.pk = d.fk;
```

# Natural Joins

Given the verbosity of straight-up JOINS, there is also a **syntax for natural joins in SQL**

```
SELECT *  
FROM A NATURAL JOIN B NATURAL JOIN C;
```

No explicit joining in the WHERE clause necessary

Requires that **primary key and foreign key attributes are named the same**



# Natural Joins

The following **does not work**:

```
SELECT E.Fname, E.Lname, D.Dname  
FROM (Employee as E) NATURAL JOIN (Department as D);
```

(PK of Department is called DNumber, but foreign key in Employee is called Dno)

Can use **aliasing** to get around this problem:

```
SELECT E.Fname, E.Lname, D.Dname  
FROM (Employee as E) NATURAL JOIN  
      (Department as D(Dname, Dno, Mssn, Msdate));
```

# Outer Joins

SQL also supports the LEFT, RIGHT, and FULL style of outer joins

```
SELECT  E.Fname, E.Lname, S.Fname, S.Lname  
FROM    EMPLOYEE AS E LEFT OUTER JOIN  
        EMPLOYEE AS S ON E.Super_ssn=S.Ssn;
```

RIGHT and FULL outer joins follow in the same style

In outer joins we are **required to use** the “On” syntax

# Unions, Intersect, Except

Finally, SQL also supports the set operators UNION, INTERSECT, and EXCEPT (difference)

**Query 4.** Make a list of all project numbers for projects that involve an employee whose last name is 'Smith', either as a worker or as a manager of the department that controls the project.

```
Q4A: ( SELECT    DISTINCT Pnumber
      FROM      PROJECT, DEPARTMENT, EMPLOYEE
      WHERE     Dnum=Dnumber AND Mgr_ssn=Ssn
              AND Lname='Smith' )

      UNION

( SELECT    DISTINCT Pnumber
  FROM      PROJECT, WORKS_ON, EMPLOYEE
  WHERE     Pnumber=Pno AND Essn=Ssn
              AND Lname='Smith' );
```



# Subqueries

Some queries require us to execute **another nested query** as part of the WHERE clause  
Typically in conjunction with **predicate logic operators**:

IN, ALL, ANY (SOME), EXISTS, UNIQUE

Find all employees with higher salary than all employees in department number 5:

```
SELECT  Lname, Fname
FROM    EMPLOYEE
WHERE   Salary > ALL ( SELECT  Salary
                        FROM    EMPLOYEE
                        WHERE   Dno=5 );
```



# Subqueries

```
Q4A:  SELECT  DISTINCT Pnumber
        FROM    PROJECT
        WHERE   Pnumber IN
                ( SELECT  Pnumber
                  FROM    PROJECT, DEPARTMENT, EMPLOYEE
                  WHERE   Dnum=Dnumber AND
                        Mgr_ssn=Ssn AND Lname='Smith' )

        OR

        Pnumber IN
        ( SELECT  Pno
          FROM    WORKS_ON, EMPLOYEE
          WHERE   Essn=Ssn AND Lname='Smith' );
```





# Subqueries can also return tuples for comparison

```
SELECT DISTINCT Essn
FROM WORKS_ON
WHERE (Pno, Hours) IN ( SELECT Pno, Hours
                        FROM WORKS_ON
                        WHERE Essn='123456789' );
```



# Using tables from the superquery in the subquery

Often is is required to reference something (tables, attributes, etc.) from the superquery in the subquery

## Needs aliases

**Query 16.** Retrieve the name of each employee who has a dependent with the same first name and is the same sex as the employee.

```
Q16:  SELECT  E.Fname, E.Lname
      FROM    EMPLOYEE AS E
      WHERE   E.Ssn IN ( SELECT  Essn
                        FROM    DEPENDENT AS D
                        WHERE   E.Fname=D.Dependent_name
                        AND E.Sex=D.Sex );
```



# Using tables from the superquery in the subquery

Note that easily leads to **very inefficient queries**

Pattern to look out for:

## **N+1 Select Problem**

(if the DBMS needs to do a **full table scan** for each tuple in the superquery)



## “WITH” Subqueries

Sometimes it is easier to just create a temporary table instead of doing a “true” subquery:

```
WITH dno_5_sales as (  
    SELECT salary from EMPLOYEE WHERE Dno = 5)  
  
SELECT Lname, Fname FROM EMPLOYEE  
WHERE Salary > ALL (dno_5_sales);
```

# Aggregate Functions

Used to summarize information from multiple tuples into a single-tuple summary

**Available aggregate functions:**

COUNT, SUM, MAX, MIN, and AVG

NULL values are discarded when aggregating



# Examples

Return a single row with summary statistics of employees:

```
SELECT SUM(Salary), MAX(Salary),  
       MIN(Salary), AVG(Salary) FROM EMPLOYEE;
```

Often used in combination with aliasing:

```
SELECT SUM(Salary) AS Total_Sal,  
       MAX(Salary) AS Highest_Sal,  
       MIN(Salary) AS Lowest_Sal,  
       AVG(Salary) AS Average_Sal  
FROM EMPLOYEE;
```

# Grouping

Common problem when aggregating:

We don't want to get the aggregate of **all** rows, but somehow group them by a specific attribute  
(**partition** relation into subsets of tuples)

Example: find the average salary **per department**

Solution: GROUP BY clause



# Grouping

Example: get the number of employees and their average salary **per department**

```
SELECT  Dno, COUNT(*), AVG (Salary)
FROM    EMPLOYEE
GROUP BY Dno;
```

Unlike other aggregate functions, this query **does not** return just one row; it returns one row **per distinct grouping value** (in that case department numbers).

Basic operating principle:

First apply the grouping (i.e., figure out how many rows the result will have)

Then apply the aggregate functions for each group



# Grouping

Important restrictions:

- the grouping attribute(s) **must** appear in the SELECT clause
- other non-aggregate attribute(s) **cannot** appear in the SELECT clause

**Wrong:**

```
SELECT  COUNT(*), AVG (Salary)
FROM    EMPLOYEE
GROUP BY Dno;
```

```
SELECT  COUNT(*), AVG (Salary), Salary
FROM    EMPLOYEE
GROUP BY Dno;
```



# Grouping and WHERE

Grouping may also be done in combination with a WHERE clause, for instance as part of a JOIN

```
SELECT  Pnumber, Pname, COUNT(*)  
FROM    PROJECT, WORKS_ON  
WHERE   Pnumber=Pno  
GROUP BY Pnumber, Pname;
```

Note that the WHERE clause is evaluated **before** the grouping happens.

# Filtering after grouping

Sometimes we need to reject (filter out) an entire group - this cannot be done with WHERE

Use HAVING:

```
SELECT  Pnumber, Pname, COUNT(*)
FROM    PROJECT, WORKS_ON
WHERE   Pnumber=Pno
GROUP BY Pnumber, Pname
HAVING  COUNT(*) > 2;
```

# Key Takeaway

## Basic SQL Query Syntax

```
SELECT <attribute and function list>  
FROM <table list>  
[ WHERE <condition> ]  
[ GROUP BY <grouping attribute(s)> ]  
[ HAVING <group condition> ]  
[ ORDER BY <attribute list> ];
```

# Key Takeaway

**Understanding basic SQL queries (SELECT, FROM, WHERE)**

**DISTINCT, ORDER BY**

**Different types of joins and their syntax**

**Subselects and WITH**

**Aggregation, Grouping, and the HAVING clause**